

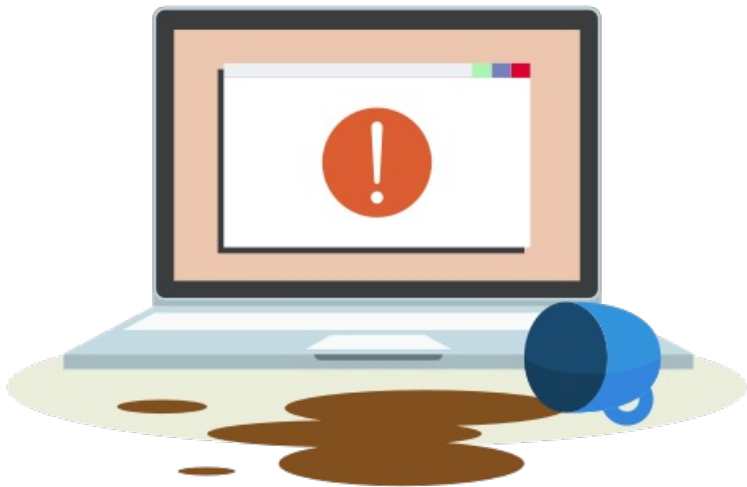
# code

**Weeks of debugging  
can save you hours of  
TLA+**

Markus A. Kuppe  
Principal Engineer – RiSE@MSR

develop & learn

Wanted to work on my slides...



# Blocking Queue

Wait:  $\{\}$

Run:  $\{p1, p2, p3, p4, c1, c2, c3\}$

Sched: p1

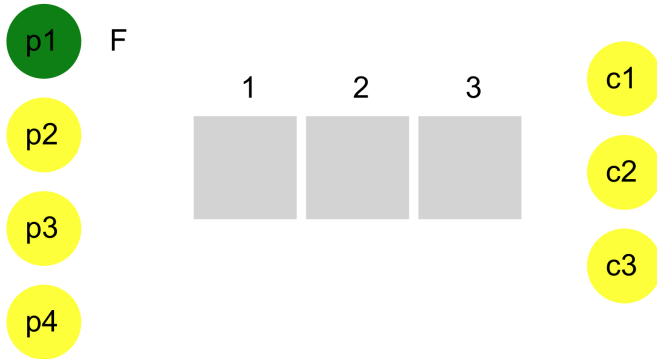


Figure: 1/47

# Blocking Queue

Wait:  $\{\}$

Run:  $\{p1, p2, p3, p4, c1, c2, c3\}$

Sched: p1

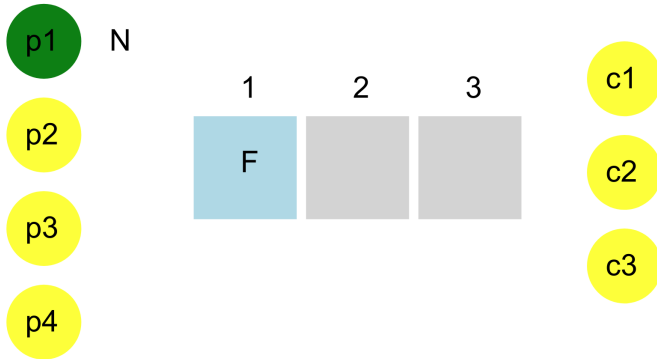


Figure: 2/47

# Blocking Queue

Wait:  $\{\}$

Run:  $\{p1, p2, p3, p4, c1, c2, c3\}$

Sched: p1

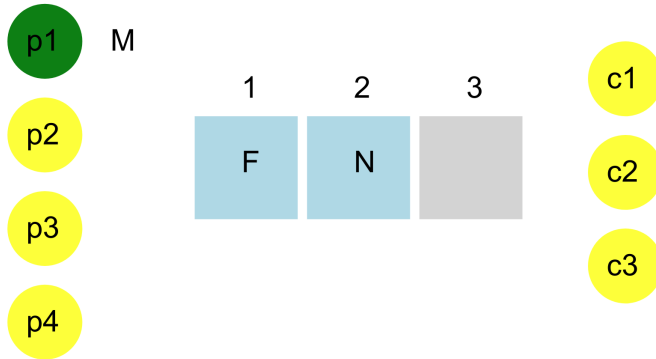


Figure: 3/47

# Blocking Queue

Wait: {p1, p2, p3, p4}

Run: {c1, c2, c3}

Sched: c1

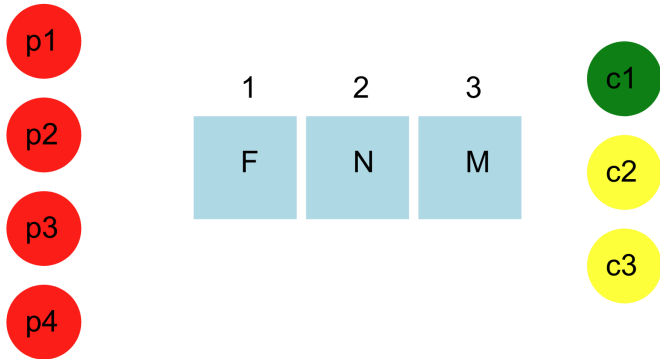


Figure: 17/47

# Blocking Queue

Wait: {p2, p3, p4}

Run: {p1, c1, c2, c3}

Sched: c1

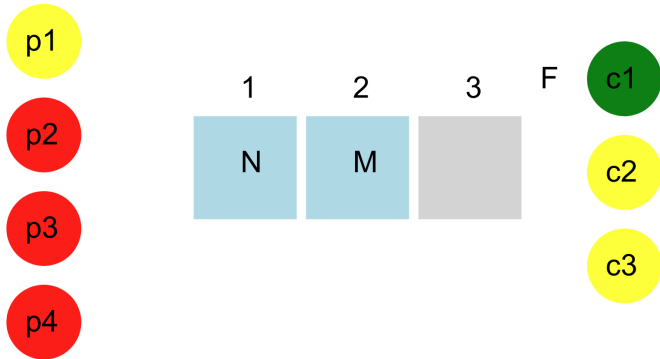


Figure: 18/47

# Blocking Queue

Wait: {p1, p2, p3, p4, c1, c2, c3}

Run: {}

Sched: p1

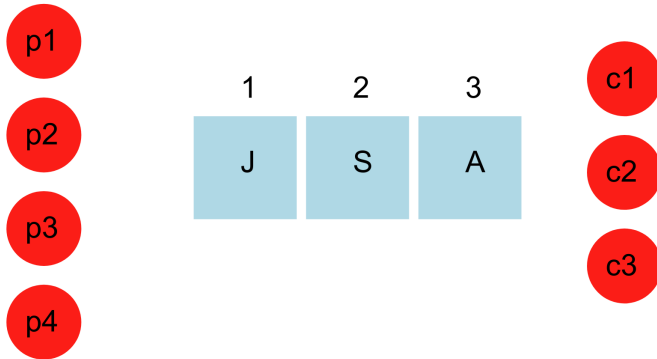


Figure: Deadlock!!!



TLA<sup>+</sup> 30.000ft above

TLA<sup>+</sup> is a specification language to design, document, and verify reactive systems.



Figure: TLA<sup>+</sup> creator

TLA<sup>+</sup> 30.000ft above

TLA<sup>+</sup> is a specification language to design, document, and verify reactive systems.

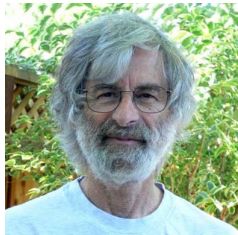


Figure: Leslie Lamport

## Specify Large Systems

- ▶ DynamoDB: scalable high-performance "no SQL" data store with cross datacenter replication and strong consistency guarantees
- ▶ First informal proofs and excessive (fault-injecting) testing

# Specify Large Systems

- ▶ DynamoDB: scalable high-performance "no SQL" data store with cross datacenter replication and strong consistency guarantees
- ▶ First informal proofs and excessive (fault-injecting) testing
- ▶ TLC found very subtle bug: shortest error trace 35 steps
- ▶ "Using TLA<sup>+</sup> in place of traditional proof writing would thus likely have improved time to market, in addition to achieving greater confidence in the system's correctness." [Newcombe, 2014, Newcombe et al., 2015]

# Specify Large Systems

- ▶ DynamoDB: scalable high-performance "no SQL" data store with cross datacenter replication and strong consistency guarantees
- ▶ First informal proofs and excessive (fault-injecting) testing
- ▶ TLC found very subtle bug: shortest error trace 35 steps
- ▶ "Using TLA<sup>+</sup> in place of traditional proof writing would thus likely have **improved time to market**, in addition to achieving **greater confidence** in the system's correctness." [Newcombe, 2014, Newcombe et al., 2015]



# Everybody Happy?

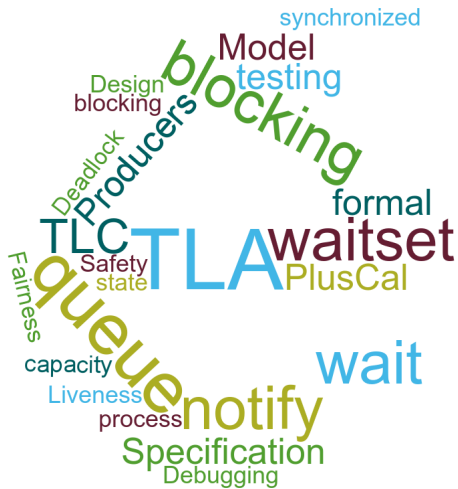
Everybody Happy?



# What Next?

- ▶ Provide Feedback on CODE!
  - ▶ <http://aka.ms/CODE-Feedback>
- ▶ Learn TLA+:
  - ▶ <http://aka.ms/TLAbq>
  - ▶ <http://aka.ms/TLA>
  - ▶ <http://aka.ms/TLAclass>
    - ▶ Sign up for a hands-on class
    - ▶ Raffle: Handsigned by LL copy of Specifying Systems book

# Q&A





## Bibliography I

Chris Newcombe. Why Amazon Chose TLA+. In David Hutchison, Takeo Kanade, Josef Kittler, Jon M. Kleinberg, Alfred Kobsa, Friedemann Mattern, John C. Mitchell, Moni Naor, Oscar Nierstrasz, C. Pandu Rangan, Bernhard Steffen, Demetri Terzopoulos, Doug Tygar, Gerhard Weikum, Yamine Ait Ameur, and Klaus-Dieter Schewe, editors, Abstract State Machines, Alloy, B, TLA, VDM, and Z, volume 8477, pages 25–39. Springer Berlin Heidelberg, Berlin, Heidelberg, 2014. ISBN 978-3-662-43651-6 978-3-662-43652-3. URL

[http://link.springer.com/10.1007/978-3-662-43652-3\\_3](http://link.springer.com/10.1007/978-3-662-43652-3_3).

Chris Newcombe, Tim Rath, Fan Zhang, Bogdan Munteanu, Marc Brooker, and Michael Deardeuff. How Amazon Web Services Uses Formal Methods. Communications of the ACM, 58(4):66–73, March 2015. ISSN 00010782. doi: 10.1145/2699417. URL

<http://dl.acm.org/citation.cfm?doid=2749359.2699417>.

“Writing is nature's way of letting you know how sloppy your thinking is.”  
Guindon



- ▶ First release of the Xbox 360
- ▶ MSR **intern** spec'ed IBM's memory coherence protocol
- ▶ **Writing** the spec revealed a subtle bug
- ▶ IBM acknowledge the bug only after several weeks
- ▶ Chips would have deadlocked after  $\sim 4$  hours of use
- ▶ Xbox Christmas launch would have been missed



## Ratio BufCapacity, Consumers, and Producers

Deadlock iff:

$$2K < |Consumers \cup Producers|$$